

Machine Learning

CMPSCI 589

Bruno C. da Silva
bsilva@cs.umass.edu

Probabilistic Classifiers

Naive Bayes (1/2)



Review: Selecting an Attribute to Test

- Decision tree to predict whether a person will play tennis

Weather	Temperature	Humidity	Windy	PlayTennis
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

- Which attribute to test first?
- Let's consider testing **Weather**

Original dataset: 9 instances "Yes"
5 instances "No"

yes yes yes yes yes yes yes yes
no no no no no

Weather

Review: Selecting an Attribute to Test

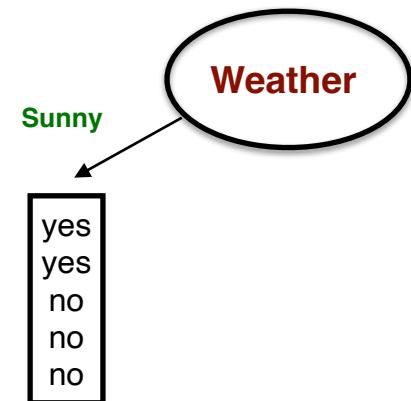
- Decision tree to predict whether a person will play tennis

Weather	Temperature	Humidity	Windy	PlayTennis
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

- Which attribute to test first?
- Let's consider testing **Weather**

Original dataset: 9 instances "Yes"
5 instances "No"

yes yes yes yes yes yes yes yes
no no no no no



Review: Selecting an Attribute to Test

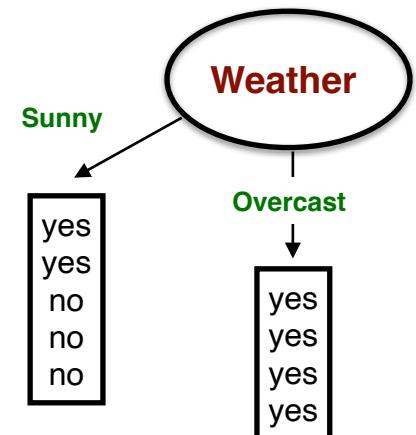
- Decision tree to predict whether a person will play tennis

Weather	Temperature	Humidity	Windy	PlayTennis
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

- Which attribute to test first?
- Let's consider testing **Weather**

Original dataset: 9 instances "Yes"
5 instances "No"

yes yes yes yes yes yes yes yes
no no no no no



Review: Selecting an Attribute to Test

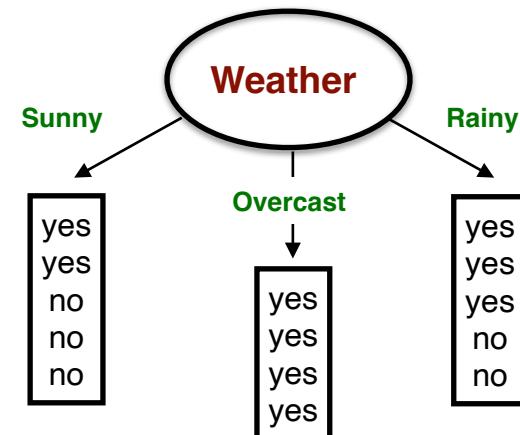
- Decision tree to predict whether a person will play tennis

Weather	Temperature	Humidity	Windy	PlayTennis
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

- Which attribute to test first?
- Let's consider testing **Weather**

Original dataset: 9 instances "Yes"
5 instances "No"

yes yes yes yes yes yes yes yes
no no no no no



Review: Selecting an Attribute to Test

- Decision tree to predict whether a person will play tennis
 - Entropy of the original dataset:
 - $I(9/14, 5/14) = -9/14 \log_2(9/14) -5/14 \log_2(5/14) = 0.940 \text{ bits}$
 - Entropy of partitions resulting from testing Weather:
 - Weather=Sunny
 - $I(2/5, 3/5) = -2/5 \log_2(2/5) -3/5 \log_2(3/5) = 0.971 \text{ bits}$
 - Weather=Overcast
 - $I(4/4, 0/4) = -4/4 \log_2(4/4) -0/4 \log_2(0/4) = 0 \text{ bits}$
 - Weather=Rainy
 - $I(3/5, 2/5) = -3/5 \log_2(3/5) -2/5 \log_2(2/5) = 0.971 \text{ bits}$
 - Average entropy of the resulting partitions
 - $(5/14) \times 0.971 + (4/14) \times 0 + (5/14) \times 0.971 = 0.693 \text{ bits}$
- Which attribute to test first?
 - Let's consider testing Weather
- Original dataset: 9 instances "Yes" 5 instances "No"
- yes yes yes yes yes yes yes yes
no no no no no
- ```
graph TD; Weather([Weather]) -- Sunny --> SunnyBox[yes
yes
no
no
no]; Weather -- Overcast --> OvercastBox[yes
yes
yes
yes]; Weather -- Rainy --> RainyBox[yes
yes
no
no];
```

The diagram shows a decision tree starting with an oval labeled "Weather". Three arrows branch out: "Sunny" leads to a box containing "yes", "yes", "no", "no", "no"; "Overcast" leads to a box containing "yes", "yes", "yes", "yes"; and "Rainy" leads to a box containing "yes", "yes", "no", "no".

# Review: Selecting an Attribute to Test

- Decision tree to predict whether a person will play tennis
- Entropy of the original dataset:
  - $I(9/14, 5/14) = -9/14 \log_2(9/14) -5/14 \log_2(5/14) = 0.940 \text{ bits}$
- Entropy of partitions resulting from testing Weather:
  - Weather=Sunny
    - $I(2/5, 3/5) = -2/5 \log_2(2/5) -3/5 \log_2(3/5) = 0.971 \text{ bits}$
  - Weather=Overcast
    - $I(4/4, 0/4) = -4/4 \log_2(4/4) -0/4 \log_2(0/4) = 0 \text{ bits}$
  - Weather=Rainy
    - $I(3/5, 2/5) = -3/5 \log_2(3/5) -2/5 \log_2(2/5) = 0.971 \text{ bits}$
- Average entropy of the resulting partitions
  - $(5/14) \times 0.971 + (4/14) \times 0 + (5/14) \times 0.971 = 0.693 \text{ bits}$

By testing the attribute **Weather**, the entropy of the classes decreased by  $0.940 - 0.693 = 0.247 \text{ bits}$



## Information Gain

# Review: Criteria for Selecting an Attribute to Test

- We have discussed one possible criterion for selecting which attribute to test
  - **Information Gain**
- Many other criteria have been proposed — each with different properties
- Intuitively:
  - A split that keeps the **same proportion of classes** in each partition is **useless**
  - A split where the **instances in each partition have the same class** is **useful!**

# Review: Criteria for Selecting an Attribute to Test

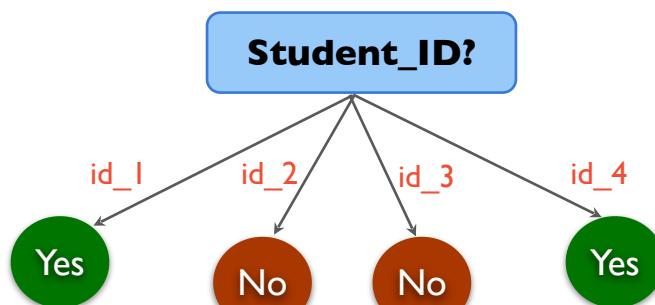
- We have discussed one possible criterion for selecting which attribute to test
  - **Information Gain**
- Many other criteria have been proposed — each with different properties
- Intuitively:
  - A split that keeps the same proportion of classes in each partition is **useless**
  - A split where the instances in each partition have the same class is **useful!**

- 
- Main criteria for selecting which attribute to test:
    - **Information Gain** - ID3 Algorithm (Quilan, 1987)
    - **Information Gain Ratio** - C4.5 Algorithm (Quilan, 1988)
    - **Gini Impurity** - CART Algorithm (Breiman, 1984)

# Review: Information Gain

- This is the criterion discussed earlier → results in a method known as **ID3**
- Intuitively, it **selects the attribute  $A$**  that **maximizes the difference between**:
  - The **entropy of the original dataset  $D$**  (before splitting it based on  $A$ )
  - The **average entropy of the resulting partitions** if we split dataset  $D$  based on  $A$
- Often results in a decision tree that is not necessarily the “simplest” one
- Intuitively, it often chooses attributes with many possible values (like **Student\_ID**, **Name**, etc)

| Student_ID | Student | Age         | Credit_Score | Will_Buy_Computer |
|------------|---------|-------------|--------------|-------------------|
| id_1       | Yes     | Young       | Regular      | Yes               |
| id_2       | Yes     | Middle Age  | Excellent    | No                |
| id_3       | No      | Young       | Excellent    | No                |
| id_4       | No      | Older Adult | Regular      | Yes               |



- Perfect split!
- With just one test, can “predict” the class perfectly
- But it is clearly overfitting (“memorizing” the dataset)

# Review: Information Gain

- This is the criterion discussed earlier → results in a method known as **ID3**
- Intuitively, it **selects the attribute  $A$**  that **maximizes the difference between**:
  - The **entropy of the original dataset  $D$**  (before splitting it based on  $A$ )
  - The **average entropy of the resulting partitions** if we split dataset  $D$  based on  $A$
- Often results in a decision tree that is not necessarily the “simplest” one
- Intuitively, it often chooses attributes with many possible values (like **Student\_ID**, **Name**, etc)

| Student_ID | Student | Age         | Credit_Score | Will_Buy_Computer |
|------------|---------|-------------|--------------|-------------------|
| id_1       | Yes     | Young       | Regular      | Yes               |
| id_2       | Yes     | Middle Age  | Excellent    | No                |
| id_3       | No      | Young       | Excellent    | No                |
| id_4       | No      | Older Adult | Regular      | Yes               |

Student\_ID?

Perfect split!

The **Information Gain Ratio** criterion, implemented by the **C4.5** algorithm, tries to mitigate this issue



With just one test, can predict the class perfectly  
But it is clearly overfitting ("memorizing" the dataset)

# Review: Criteria for Selecting an Attribute to Test

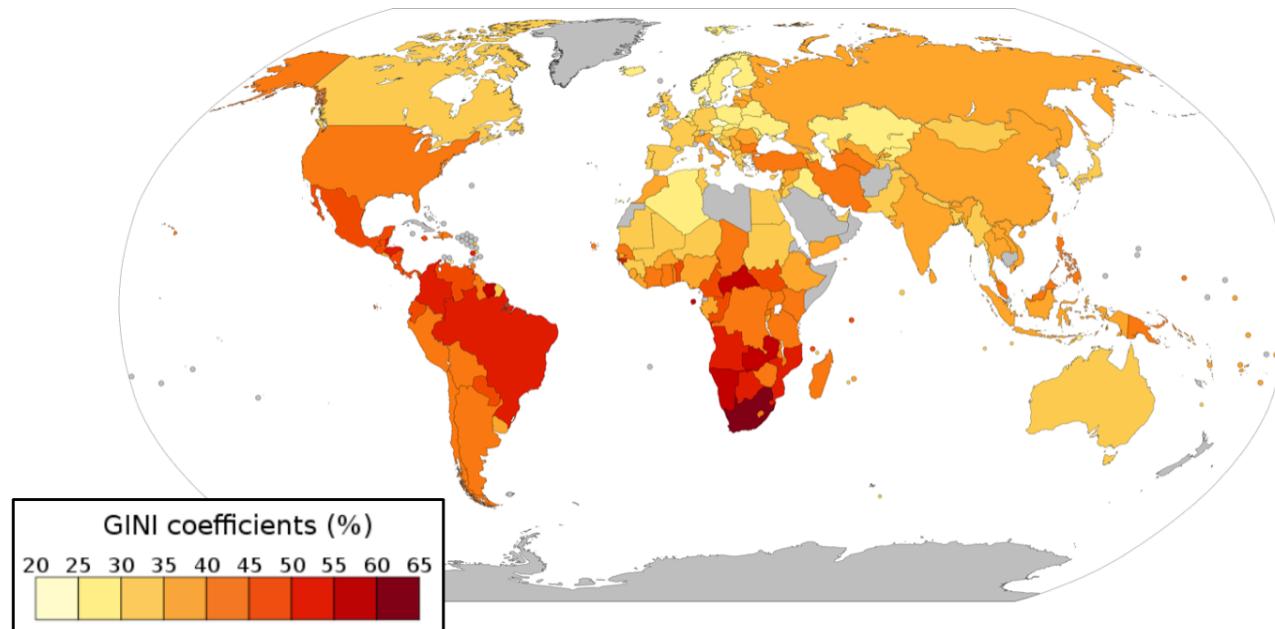
- We have discussed one possible criterion for selecting which attribute to test
  - **Information Gain**
- Many other criteria have been proposed — each with different properties
- Intuitively:
  - A split that keeps the same proportion of classes in each partition is **useless**
  - A split where the instances in each partition have the same class is **useful!**

- Main criteria for selecting which attribute to test:
  - **Information Gain** - ID3 Algorithm (Quilan, 1987)
  - **Information Gain Ratio** - C4.5 Algorithm (Quilan, 1988)
  - **Gini Impurity** - CART Algorithm (Breiman, 1984)



# Gini Criterion

- Originally proposed to quantify how uneven income is across a population



- Gini coefficient → how uneven income/wealth distribution across a population is
  - $\text{Gini} = 1 \rightarrow$  very uneven income/wealth distribution across a population
  - $\text{Gini} = 0 \rightarrow$  very even income/wealth distribution across a population

# Gini Criterion

- Gini coefficient → how uneven income/wealth distribution across a population is
- In the context of decision trees
  - how uneven (or non-homogeneous) are the classes after a split

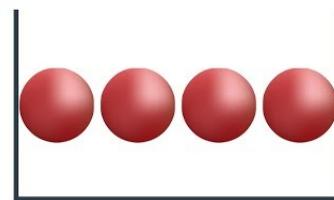
# Gini Criterion

- Gini coefficient → how uneven income/wealth distribution across a population is
- In the context of decision trees
  - how uneven (or non-homogeneous) are the classes after a split

**Let's suppose we test Age, and the instances associated with Age=Young look like this**

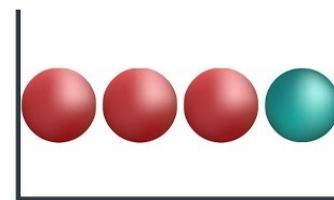
 Will repay loan

 Will not repay loan



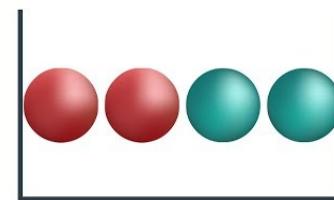
**Even**

$$\begin{aligned} \Pr(\text{Red}) &= 1 \\ \Pr(\text{Teal}) &= 0 \end{aligned}$$



**“Medium”**

$$\begin{aligned} \Pr(\text{Red}) &= 3/4 \\ \Pr(\text{Teal}) &= 1/4 \end{aligned}$$



**Uneven**

$$\begin{aligned} \Pr(\text{Red}) &= 2/4 \\ \Pr(\text{Teal}) &= 2/4 \end{aligned}$$

# Gini Criterion

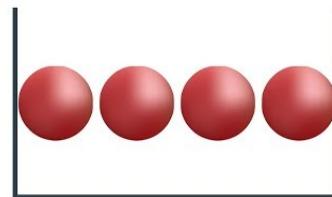
- Gini coefficient → how uneven income/wealth distribution across a population is
- In the context of decision trees
  - how uneven (or non-homogeneous) are the classes after a split

$$\text{Gini}(D) = 1 - (\Pr(\text{Red})^2 + \Pr(\text{Green})^2)$$

Let's suppose we test Age, and the instances associated with Age=Young look like this

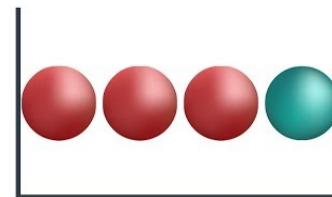
Red circle: Will repay loan

Green circle: Will not repay loan



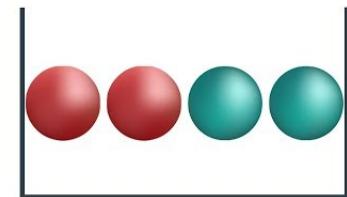
Even

$$\begin{aligned}\Pr(\text{Red}) &= 1 \\ \Pr(\text{Green}) &= 0\end{aligned}$$



“Medium”

$$\begin{aligned}\Pr(\text{Red}) &= 3/4 \\ \Pr(\text{Green}) &= 1/4\end{aligned}$$



Uneven

$$\begin{aligned}\Pr(\text{Red}) &= 2/4 \\ \Pr(\text{Green}) &= 2/4\end{aligned}$$

Gini  $\Rightarrow$

$$\begin{aligned}1 - (1^2 + 0^2) \\ = 0\end{aligned}$$

$$\begin{aligned}1 - ((3/4)^2 + (1/4)^2) \\ = 0.375\end{aligned}$$

$$\begin{aligned}1 - ((2/4)^2 + (2/4)^2) \\ = 0.5\end{aligned}$$

# Gini Criterion

- Gini coefficient → how uneven income/wealth distribution across a population is
- In the context of decision trees
  - how uneven (or non-homogeneous) are the classes after a split

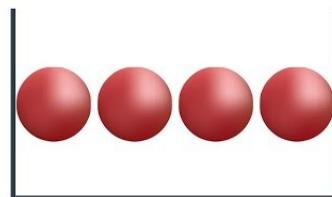
$$\text{Gini}(D) = 1 - (\Pr(\text{Red})^2 + \Pr(\text{Green})^2)$$

Let's suppose we test Age, and the instances associated with Age=Young look like this

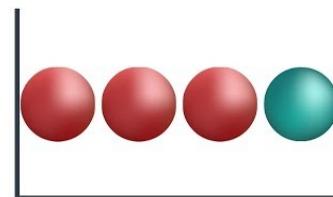


Will repay loan

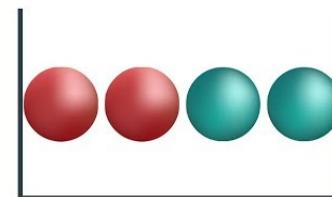
Will not repay loan



Even



“Medium”



Uneven

Gini  $\Rightarrow$

$$\begin{aligned} 1 - (1^2 + 0^2) \\ = 0 \end{aligned}$$

$$\begin{aligned} 1 - ((3/4)^2 + (1/4)^2) \\ = 0.375 \end{aligned}$$

$$\begin{aligned} 1 - ((2/4)^2 + (2/4)^2) \\ = 0.5 \end{aligned}$$



more homogenous partition → ideal result of a split  
(smaller value of the Gini coefficient)

# Gini Criterion

- In the context of decision trees
  - how uneven (or non-homogeneous) are the classes after a split

$$\text{Gini}(D) = 1 - (\Pr(\text{Red})^2 + \Pr(\text{Green})^2)$$

- More generally, if there are  $m$  classes in a dataset  $D$

$$\text{Gini}(D) = 1 - \left( \sum_{i=1}^m (p_i)^2 \right)$$

where  $p_i$  be the probability that the label/class  $i$  occurs in instances in a dataset  $D$

# Gini Criterion

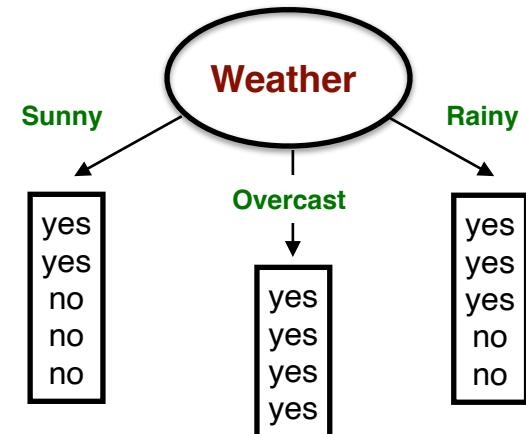
- Decision tree to predict whether a person will play tennis

| Weather  | Temperature | Humidity | Windy | PlayTennis |
|----------|-------------|----------|-------|------------|
| Sunny    | Hot         | High     | False | No         |
| Sunny    | Hot         | High     | True  | No         |
| Overcast | Hot         | High     | False | Yes        |
| Rainy    | Mild        | High     | False | Yes        |
| Rainy    | Cool        | Normal   | False | Yes        |
| Rainy    | Cool        | Normal   | True  | No         |
| Overcast | Cool        | Normal   | True  | Yes        |
| Sunny    | Mild        | High     | False | No         |
| Sunny    | Cool        | Normal   | False | Yes        |
| Rainy    | Mild        | Normal   | False | Yes        |
| Sunny    | Mild        | Normal   | True  | Yes        |
| Overcast | Mild        | High     | True  | Yes        |
| Overcast | Hot         | Normal   | False | Yes        |
| Rainy    | Mild        | High     | True  | No         |

- Let's consider testing Weather

Original dataset: 9 instances "Yes"  
5 instances "No"

yes yes yes yes yes yes yes yes  
no no no no no



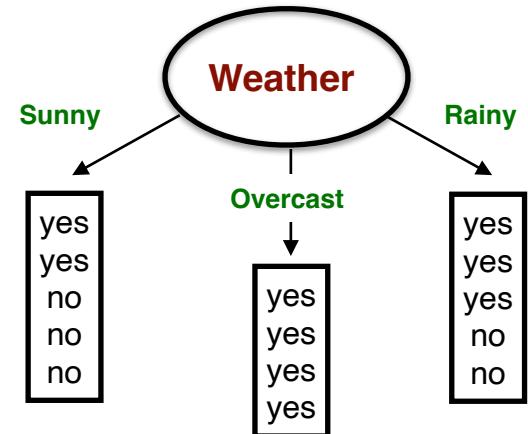
# Gini Criterion

- Decision tree to predict whether a person will play tennis
- Gini coefficient of the original dataset:
  - $\text{Gini}(9/14, 5/14) = 1 - ((9/14)^2 + (5/14)^2) = 0.459$
- Gini coeff. of partitions resulting from testing Weather:
  - Weather=Sunny
    - $\text{Gini}_{\text{Sunny}}(2/5, 3/5) = 1 - ((2/5)^2 + (3/5)^2) = 0.48$
  - Weather=Overcast
    - $\text{Gini}_{\text{Overcast}}(4/4, 0/4) = 1 - ((4/4)^2 + (0/4)^2) = 0$
  - Weather=Rainy
    - $\text{Gini}_{\text{Rainy}}(3/5, 2/5) = 1 - ((3/5)^2 + (2/5)^2) = 0.48$
- Average Gini coefficient of the resulting partitions
  - $(5/14) \times 0.48 + (4/14) \times 0 + (5/14) \times 0.48 = 0.3428$

- Let's consider testing Weather

Original dataset: 9 instances "Yes"  
5 instances "No"

yes yes yes yes yes yes yes yes yes  
no no no no no



# Gini Criterion

- Decision tree to predict whether a person will play tennis
- Gini coefficient of the original dataset:
  - $\text{Gini}(9/14, 5/14) = 1 - ((9/14)^2 + (5/14)^2) = 0.459$
- Gini coeff. of partitions resulting from testing Weather:
  - Weather=Sunny
    - $\text{Gini}_{\text{Sunny}}(2/5, 3/5) = 1 - ((2/5)^2 + (3/5)^2) = 0.48$
  - Weather=Overcast
    - $\text{Gini}_{\text{Overcast}}(4/4, 0/4) = 1 - ((4/4)^2 + (0/4)^2) = 0$
  - Weather=Rainy
    - $\text{Gini}_{\text{Rainy}}(3/5, 2/5) = 1 - ((3/5)^2 + (2/5)^2) = 0.48$
- Average Gini coefficient of the resulting partitions
  - $(5/14) \times 0.48 + (4/14) \times 0 + (5/14) \times 0.48 = 0.3428$

Testing the attribute **Weather**:

$$\text{Gini}(\text{Weather}) = 0.3428$$



- Now proceed similarly as when selecting attributes via Information Gain...
- Compute Gini coefficient of each candidate attribute
- Split dataset using the attribute with the **lowest Gini coefficient**



# Criteria for Selecting an Attribute to Test

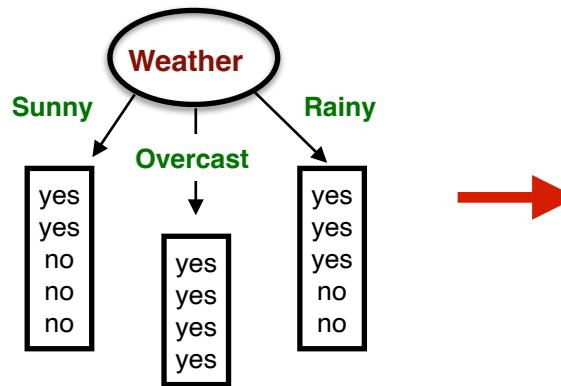
- Main criteria for selecting which attribute to test:
  - **Information Gain** - ID3 Algorithm (Quilan, 1987)
  - **Information Gain Ratio** - C4.5 Algorithm (Quilan, 1988)
  - **Gini Impurity** - CART Algorithm (Breiman, 1984)

- **Empirically:**

- **Information Gain Ratio** is almost always better than **Information Gain**
  - in terms of **predictive power** and **complexity of the resulting decision trees**
- However, in practice
  - which criterion will work best depends heavily on the application
  - should test them all and compare the resulting performances

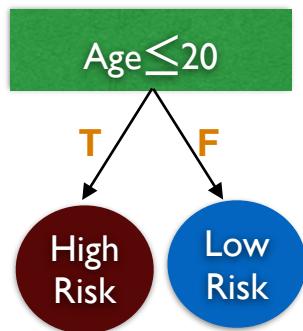
# Dealing with Numerical Attributes

- So far we have studied how to select which **categorical attribute** to split



One branch per possible value of the attribute

- How do we decide a splitting point/value in case of **numerical attributes**?



Consider deciding how to split the attribute **Age**

Pick a threshold value,  $V$

Generate two branches/disjoint partitions:

- one partition with instances s.t.  $\text{Age} \leq V$
- one partition with instances s.t.  $\text{Age} > V$

# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?
  - one partition with instances s.t.  $\text{Age} \leq V$
  - one partition with instances s.t.  $\text{Age} > V$

## 1) Sort the instances according to the value of the attribute

| Name   | Age | Gender | TrafficTicket | Class:<br>High-Risk Driver |
|--------|-----|--------|---------------|----------------------------|
| John   | 43  | M      | Yes           | High Risk                  |
| Peter  | 18  | M      | No            | High Risk                  |
| Anna   | 35  | F      | No            | Low Risk                   |
| Paula  | 19  | F      | No            | High Risk                  |
| Mark   | 90  | M      | Yes           | High Risk                  |
| Marisa | 21  | F      | Yes           | High Risk                  |
| Bob    | 30  | M      | No            | Low Risk                   |

# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?
  - one partition with instances s.t.  $\text{Age} \leq V$
  - one partition with instances s.t.  $\text{Age} > V$

## 1) Sort the instances according to the value of the attribute

| Name   | Age | Gender | TrafficTicket | Class:<br>High-Risk Driver |
|--------|-----|--------|---------------|----------------------------|
| Peter  | 18  | M      | No            | High Risk                  |
| Paula  | 19  | F      | No            | High Risk                  |
| Marisa | 21  | F      | Yes           | High Risk                  |
| Bob    | 30  | M      | No            | Low Risk                   |
| Anna   | 35  | F      | No            | Low Risk                   |
| John   | 43  | M      | Yes           | High Risk                  |
| Mark   | 90  | M      | Yes           | High Risk                  |



# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?

- one partition with instances s.t.  $\text{Age} \leq V$
- one partition with instances s.t.  $\text{Age} > V$

1) Sort the instances according to the value of the attribute

2) Evaluate splits done using as threshold the mean values between consecutive Ages

| Name   | Age | Gender | TrafficTicket | Class:<br>High-Risk Driver |
|--------|-----|--------|---------------|----------------------------|
| Peter  | 18  | M      | No            | High Risk                  |
| Paula  | 19  | F      | No            | High Risk                  |
| Marisa | 21  | F      | Yes           | High Risk                  |
| Bob    | 30  | M      | No            | Low Risk                   |
| Anna   | 35  | F      | No            | Low Risk                   |
| John   | 43  | M      | Yes           | High Risk                  |
| Mark   | 90  | M      | Yes           | High Risk                  |

# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?

- one partition with instances s.t.  $\text{Age} \leq V$
- one partition with instances s.t.  $\text{Age} > V$

1) Sort the instances according to the value of the attribute

2) Evaluate splits done using as threshold the mean values between consecutive Ages

| Name   | Age | Gender | TrafficTicket | Class:<br>High-Risk Driver |               |
|--------|-----|--------|---------------|----------------------------|---------------|
| Peter  | 18  | M      | No            | High Risk                  | $\Rightarrow$ |
| Paula  | 19  | F      | No            | High Risk                  |               |
| Marisa | 21  | F      | Yes           | High Risk                  |               |
| Bob    | 30  | M      | No            | Low Risk                   |               |
| Anna   | 35  | F      | No            | Low Risk                   |               |
| John   | 43  | M      | Yes           | High Risk                  |               |
| Mark   | 90  | M      | Yes           | High Risk                  | $\Rightarrow$ |

A red arrow points downwards from the 'Age' column to the data row for Peter.

A green box on the right contains the text  $\text{Age} \leq 18.5$ .

# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?

- one partition with instances s.t.  $\text{Age} \leq V$
- one partition with instances s.t.  $\text{Age} > V$

1) Sort the instances according to the value of the attribute

2) Evaluate splits done using as threshold the mean values between consecutive Ages

| Name   | Age | Gender | TrafficTicket | Class:<br>High-Risk Driver |                        |
|--------|-----|--------|---------------|----------------------------|------------------------|
| Peter  | 18  | M      | No            | High Risk                  | $\text{Age} \leq 18.5$ |
| Paula  | 19  | F      | No            | High Risk                  | $\text{Age} \leq 20$   |
| Marisa | 21  | F      | Yes           | High Risk                  |                        |
| Bob    | 30  | M      | No            | Low Risk                   |                        |
| Anna   | 35  | F      | No            | Low Risk                   |                        |
| John   | 43  | M      | Yes           | High Risk                  |                        |
| Mark   | 90  | M      | Yes           | High Risk                  |                        |

A red arrow points downwards from the 'Age' column of the table towards the bottom of the page.

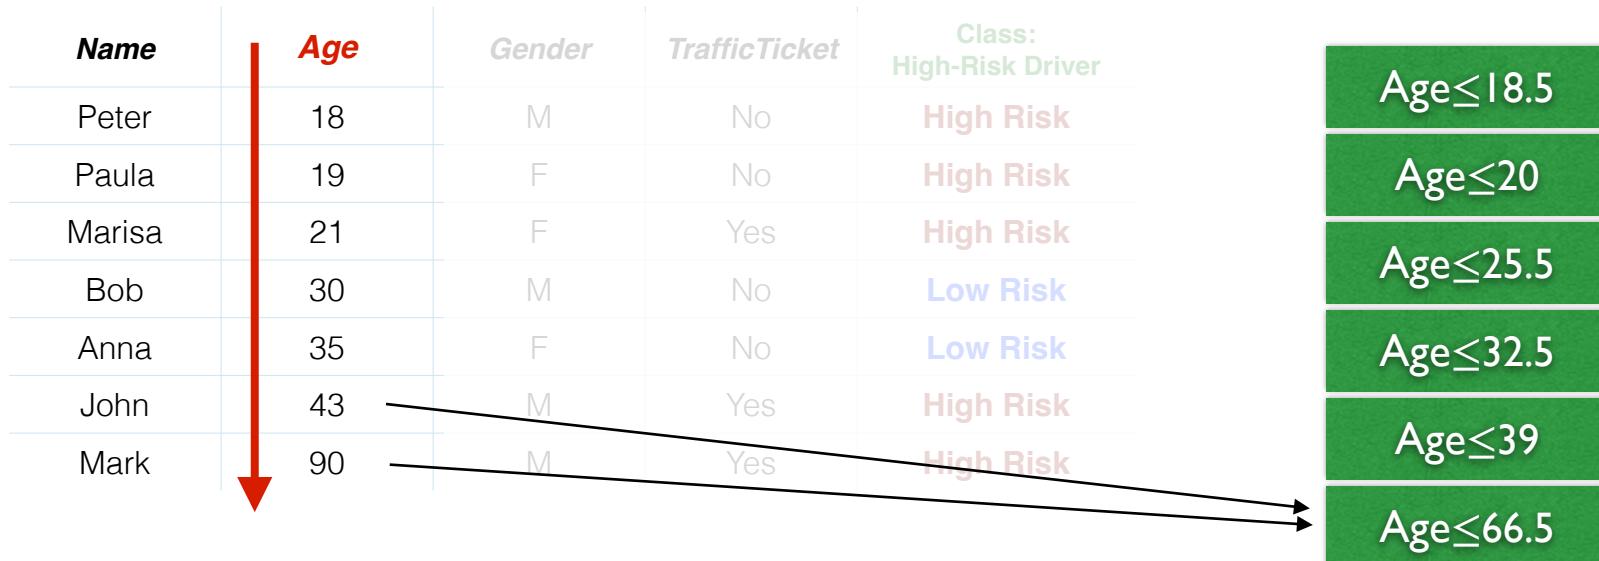
# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of numerical attributes?

- one partition with instances s.t.  $\text{Age} \leq V$
- one partition with instances s.t.  $\text{Age} > V$

1) Sort the instances according to the value of the attribute

2) Evaluate splits done using as threshold the mean values between consecutive Ages



# Dealing with Numerical Attributes

- How do we decide a splitting point/value in case of **numerical attributes**?
    - one partition with instances s.t.  $\text{Age} \leq V$
    - one partition with instances s.t.  $\text{Age} > V$
- 1) Sort the instances according to the value of the attribute
- 2) Evaluate splits done using as threshold the mean values between consecutive Ages

Age  $\leq 18.5$

Age  $\leq 20$

Age  $\leq 25.5$

Age  $\leq 32.5$

Age  $\leq 39$

Age  $\leq 66.5$

- 3) Pick the split threshold that maximizes the criterion of interest (*Info. Gain, Gini, etc.*)

- It has been shown that, for most commonly-used splitting criteria
  - testing only thresholds that correspond to such mean values is sufficient

# Decision Trees: Pros and Cons

- **Pros:**

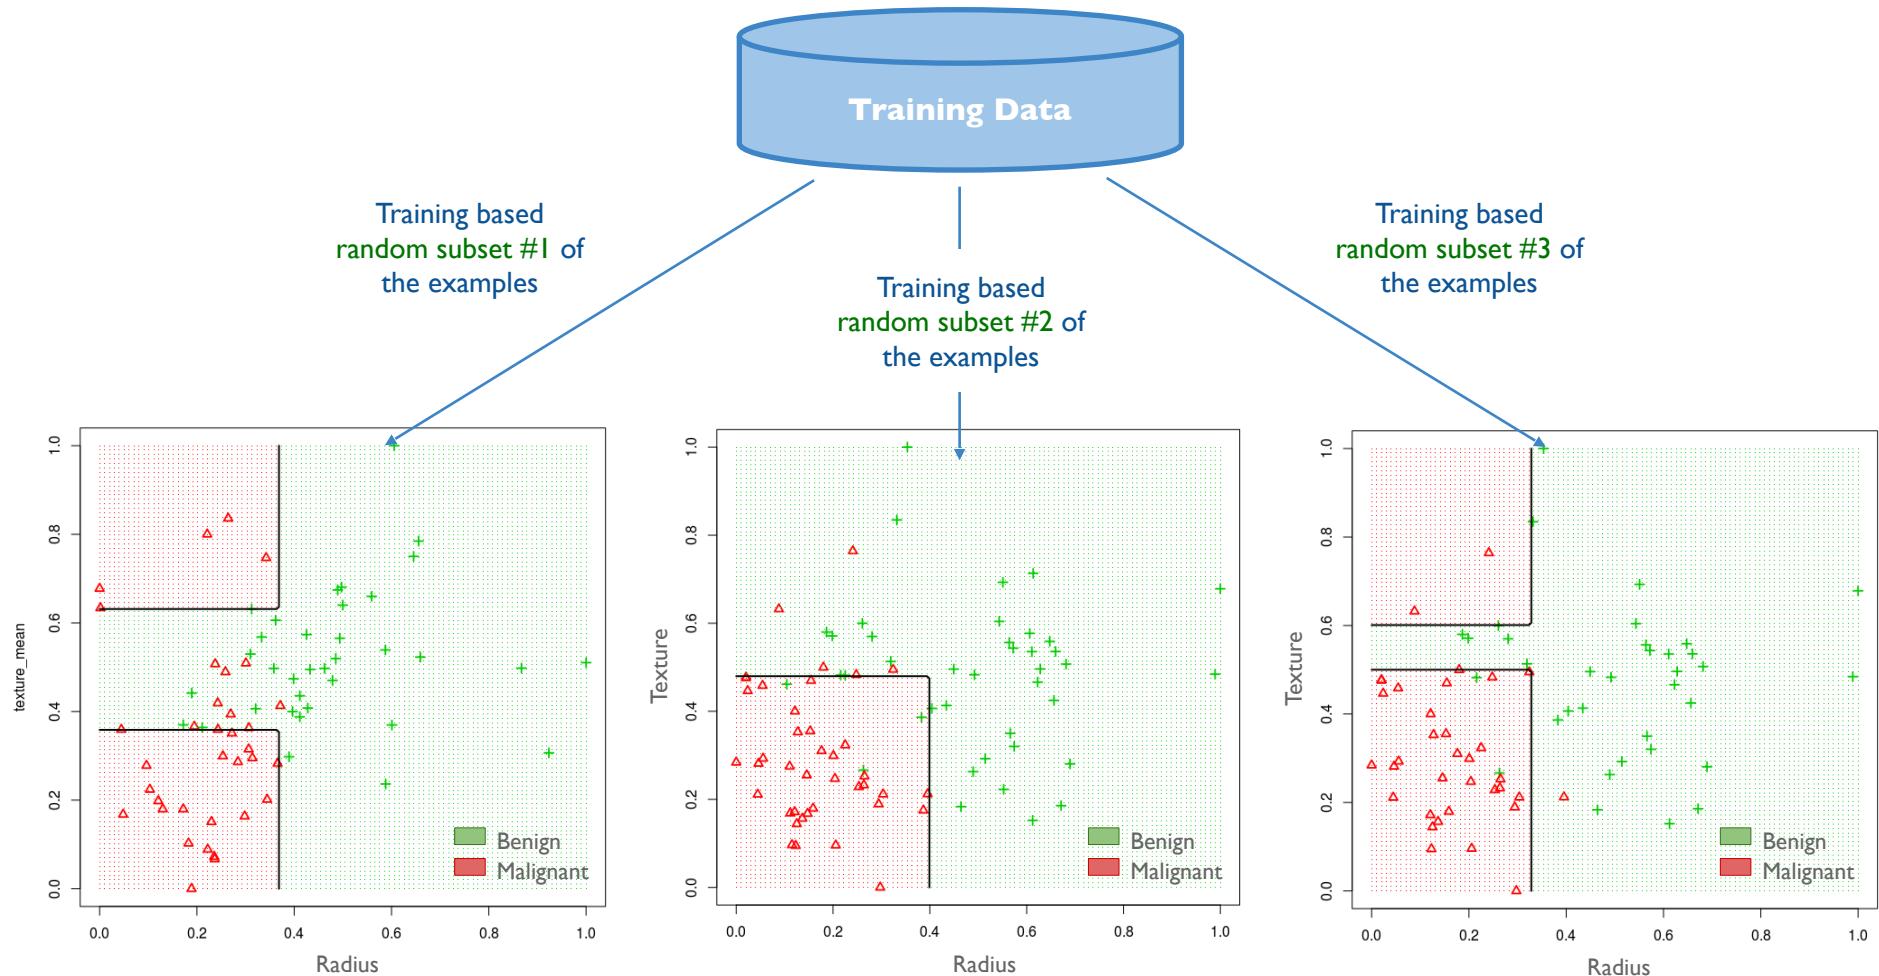
- Simple for humans to understand and interpret
- Handles both numerical and categorical attributes
- Requires little data preparation (e.g., *no need to normalize attributes*)
- Performs well with large datasets
- “Automatically” ignores irrelevant attributes not useful to predict the class/label

- **Cons:**

- Non-robust: small variations in the dataset can generate completely different trees

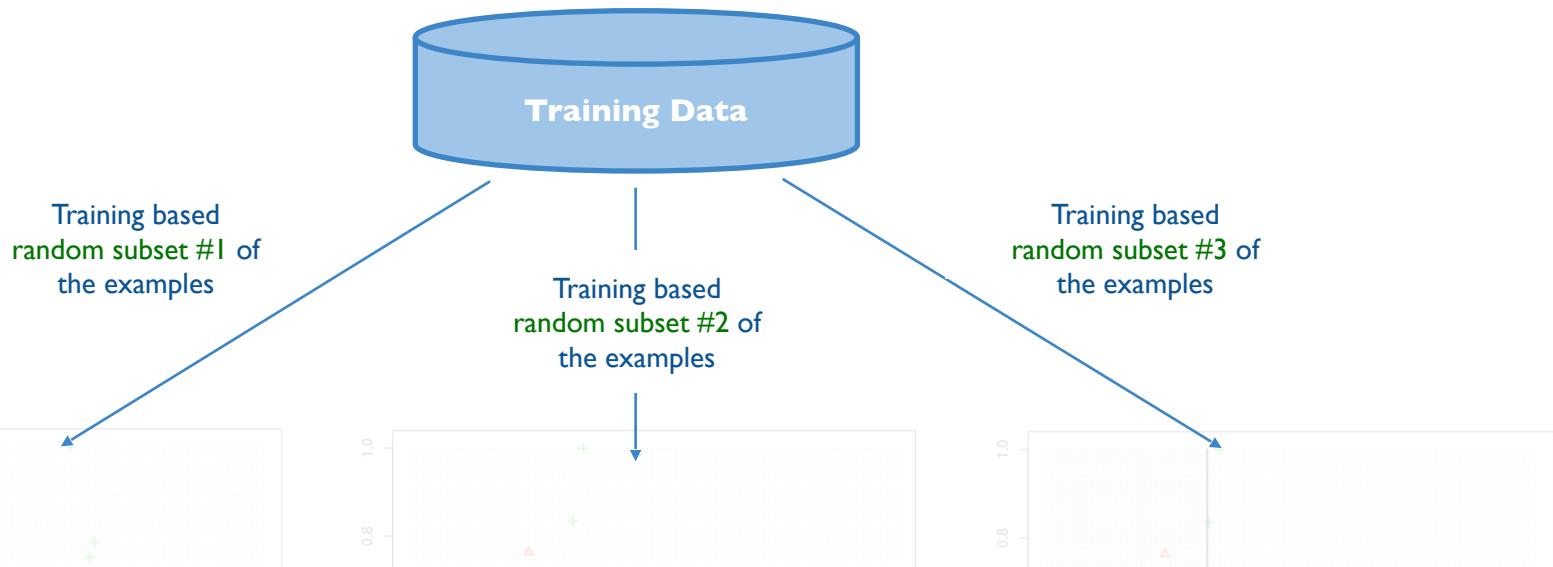
# Decision Trees Can be Non-Robust

Decision boundary can be **heavily influenced** by **small changes** to the training data



# Decision Trees Can be Non-Robust

Decision boundary can be **heavily influenced** by **small changes** to the training data



Tackling this issue is the motivation to develop the Random Forests algorithm (will be studied later)



# Decision Trees: Pros and Cons

- **Pros:**

- Simple for humans to understand and interpret
- Handles both numerical and categorical attributes
- Requires little data preparation (e.g., *no need to normalize attributes*)
- Performs well with large datasets
- “Automatically” ignores irrelevant attributes not useful to predict the class/label

- **Cons:**

- Non-robust: small variations in the dataset can generate completely different trees
- Often generate overly-complicated trees that overfit to training data
  - i.e., that do not generalize well (make correct predictions) to new instances
- Although it is possible to deal with numerical attributes, it is time-consuming
  - estimates suggest that processing them takes ~70% of execution time (Catlett, 1991)

# Naive Bayes Algorithm

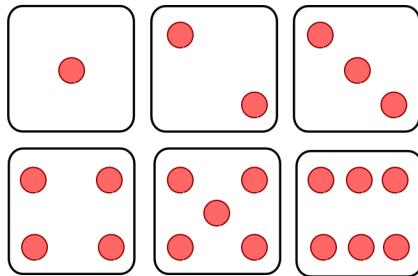


but before that...

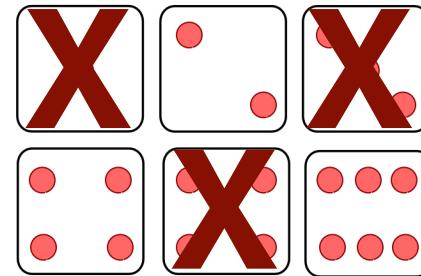
# Probability Theory - Review



# Review: Probability Theory



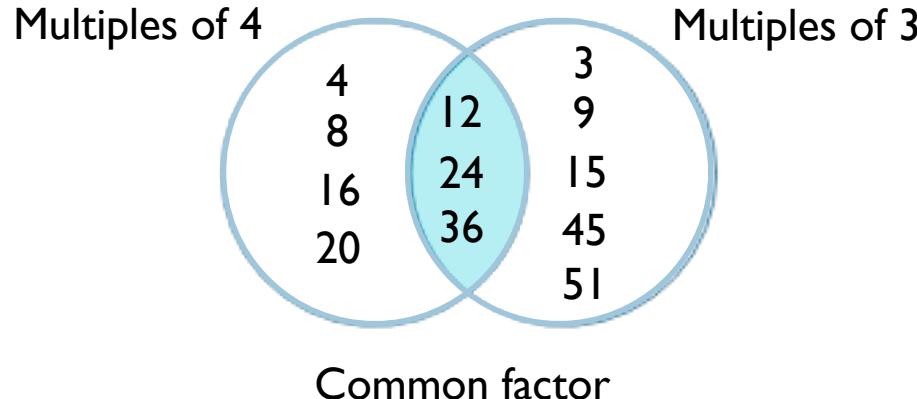
Die: probability of rolling a 2 = **1/6**



Die: probability of rolling a 2  
given that it's a special die  
that only produces even numbers

= **1/3**

# Review: Probability Theory

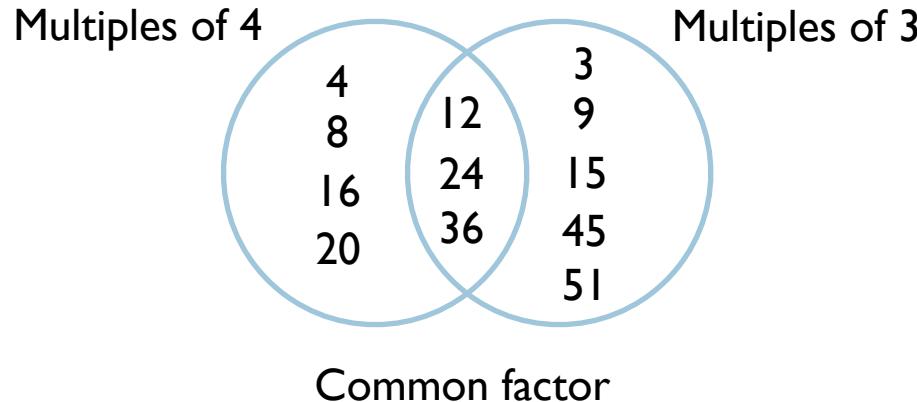


$$\Pr(\text{Multiple of 4}) = 7/12$$

$$\Pr(\text{Multiple of 3}) = 8/12$$

$$\Pr(\text{Multiple of 4 and Multiple of 3}) = 3/12$$

# Review: Probability Theory



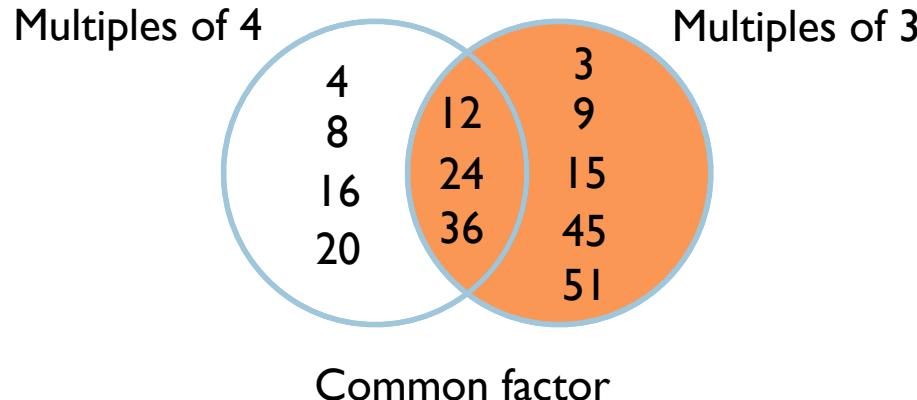
$$\Pr(\text{Multiple of 4}) = 7/12$$

$$\Pr(\text{Multiple of 3}) = 8/12$$

$$\Pr(\text{Multiple of 4 and Multiple of 3}) = 3/12$$

$$\Pr(\text{Multiple of 4} \mid \text{Multiple of 3})$$

# Review: Probability Theory



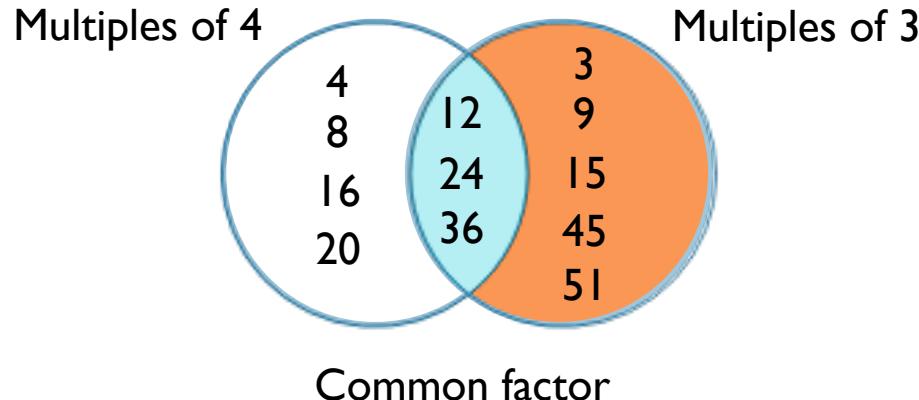
$$\Pr(\text{Multiple of 4}) = 7/12$$

$$\Pr(\text{Multiple of 3}) = 8/12$$

$$\Pr(\text{Multiple of 4 and Multiple of 3}) = 3/12$$

$$\Pr(\text{Multiple of 4} \mid \text{Multiple of 3})$$

# Review: Probability Theory



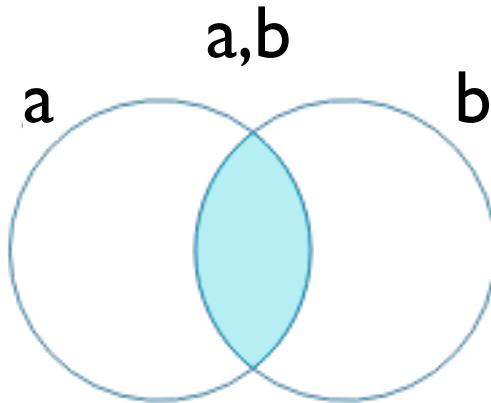
$$\Pr(\text{Multiple of 4}) = 7/12$$

$$\Pr(\text{Multiple of 3}) = 8/12$$

$$\Pr(\text{Multiple of 4 and Multiple of 3}) = 3/12$$

$$\Pr(\text{Multiple of 4} \mid \text{Multiple of 3}) = \frac{\Pr(\text{Multiple of 4 and Multiple of 3})}{\Pr(\text{Multiple of 3})}$$

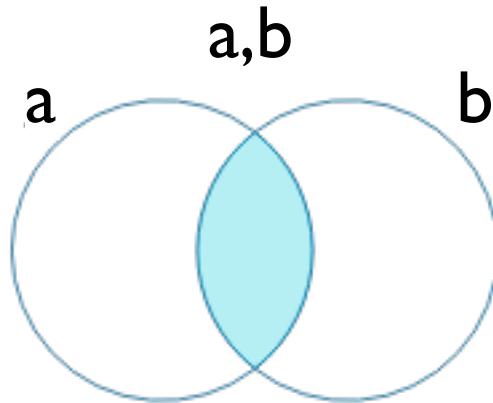
# Review: Probability Theory



$$\Pr(\text{Multiple of 4} \mid \text{Multiple of 3}) = \frac{\Pr(\text{Multiple of 4 and Multiple of 3})}{\Pr(\text{Multiple of 3})}$$

$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

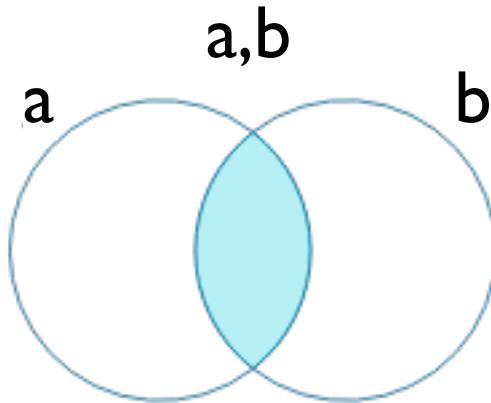
# Review: Probability Theory



$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

# Review: Probability Theory



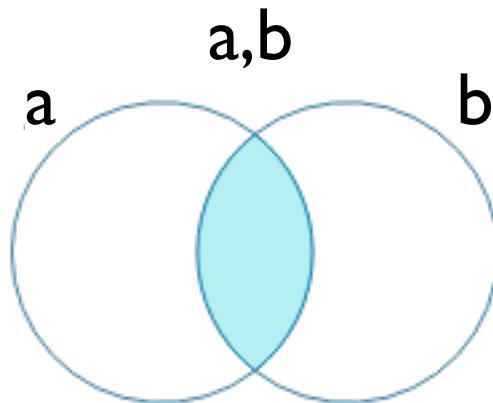
$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

$$\Pr(b \mid a) \Pr(a) = \Pr(b, a)$$

# Review: Probability Theory



$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

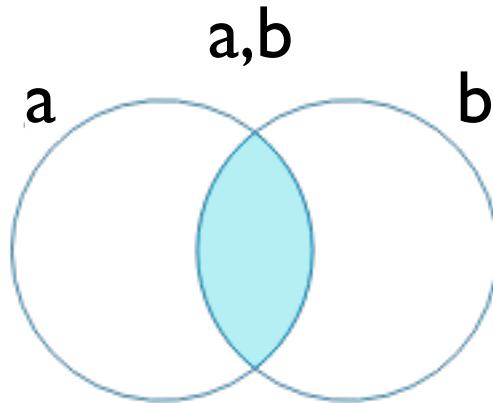
$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

$$\Pr(b \mid a) \Pr(a) = \Pr(b, a)$$

$$\Pr(a, b) = \Pr(b, a)$$



# Review: Probability Theory



$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

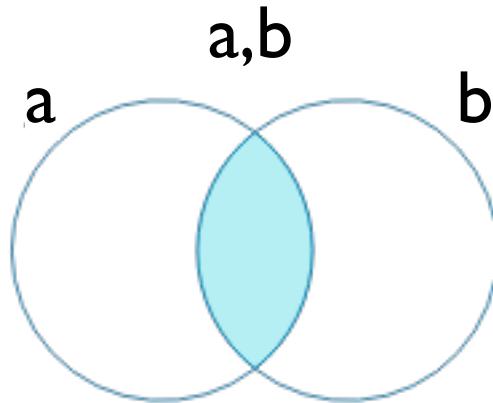
$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

$$\Pr(b \mid a) \Pr(a) = \Pr(a, b)$$

$$\Pr(a, b) = \Pr(b, a)$$



# Review: Probability Theory



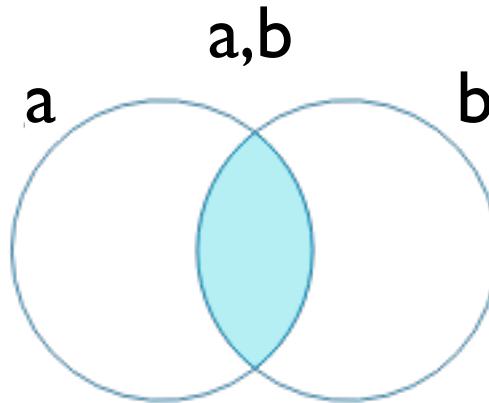
$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

$$\Pr(b \mid a) \Pr(a) = \Pr(a, b)$$

# Review: Probability Theory



$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

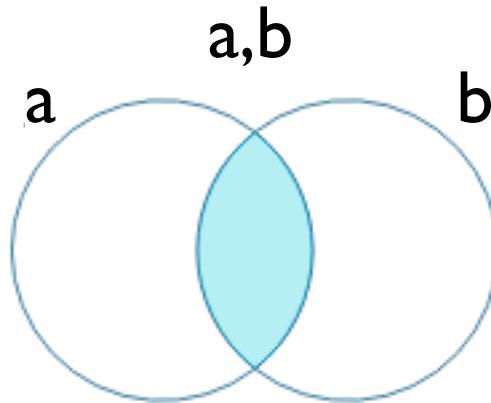
$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

$$\Pr(a \mid b) \Pr(b) = \Pr(a, b)$$

$$\Pr(b \mid a) \Pr(a) = \Pr(a, b)$$

$$\Pr(a \mid b) \Pr(b) = \Pr(b \mid a) \Pr(a)$$

# Review: Probability Theory



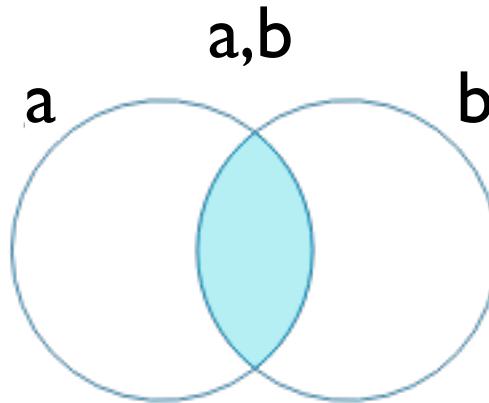
$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

$$\Pr(a \mid b) \Pr(b) = \Pr(b \mid a) \Pr(a)$$

$$\Pr(a \mid b) = \frac{\Pr(b \mid a) \Pr(a)}{\Pr(b)}$$

# Review: Probability Theory



$$\Pr(a \mid b) = \frac{\Pr(a, b)}{\Pr(b)}$$

$$\Pr(b \mid a) = \frac{\Pr(b, a)}{\Pr(a)}$$

$$\Pr(a \mid b) = \frac{\Pr(b \mid a) \Pr(a)}{\Pr(b)}$$

## Bayes Theorem



# Machine Learning

## CMPSCI 589

**Bruno C. da Silva**  
[bsilva@cs.umass.edu](mailto:bsilva@cs.umass.edu)

## Probabilistic Classifiers

### Naive Bayes (1/2)

